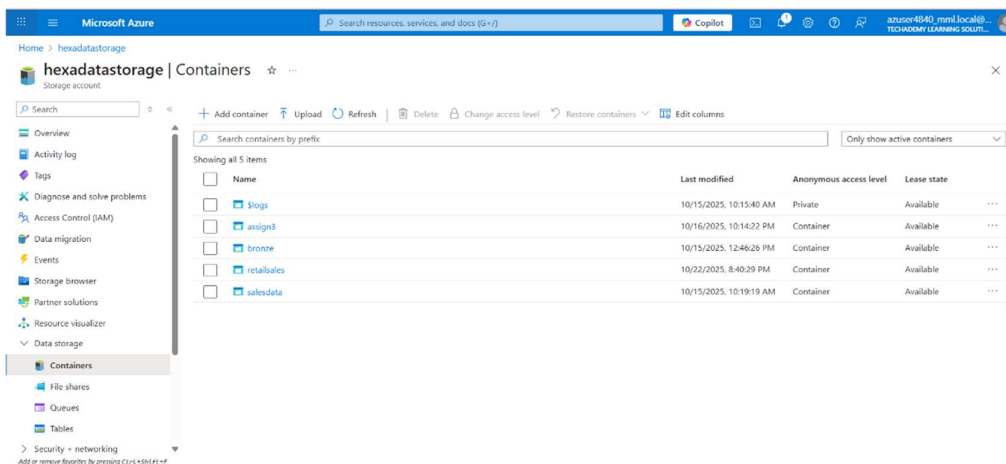
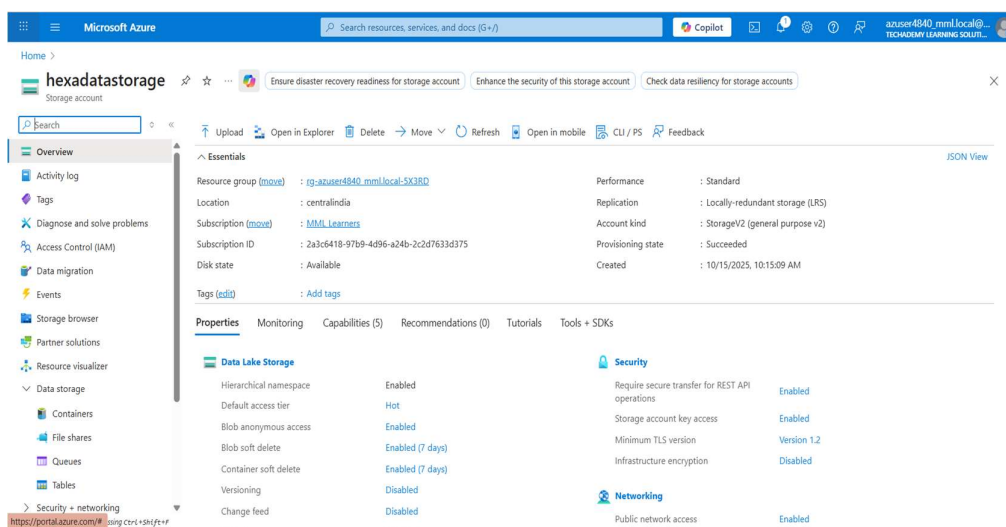
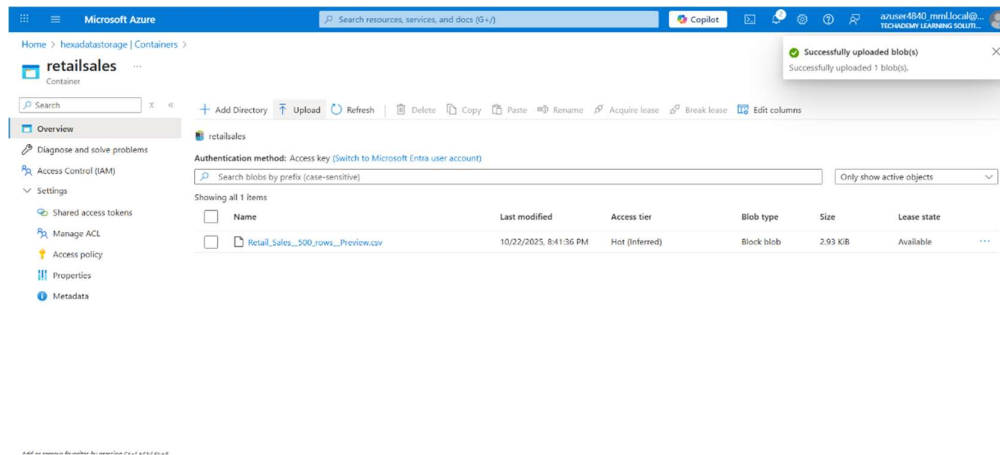


Case Study: Ingesting and Modeling Retail Sales Data from Azure to Snowflake using Snowpark, Visualized in Power BI

Step 1: Upload Data to Azure Storage

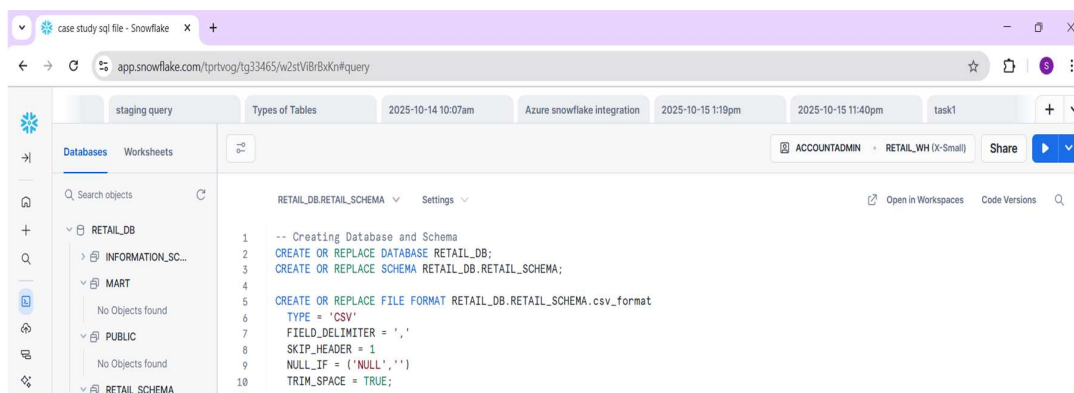
- Created an Azure Storage Account named **hexadatastorage**.
- Created a container named **retailsales** inside the storage account.
- Uploaded the file **Retail_Sales__500_rows__Preview.csv** into the container.
- Verified the file upload in the Azure Portal.





Step 2: Ingest Data into Snowflake using Snowpark

- Created database **RETAIL_DB** and schema **RETAIL_SCHEMA** in Snowflake.



- Connected to Snowflake from Snowpark.



- Created an internal stage to store incoming files.
- Uploaded the CSV file from local/Azure to the Snowflake stage.
- Loaded the CSV from stage into a **Snowpark DataFrame**.
- Verified that the data loaded successfully by displaying sample records.

```

# Put file to stage
session.file.put("Retail_Sales.csv", "@AZURE_STAGE", auto_compress=False)

# Read file from stage using Snowpark
df = session.read.options({
    "field_delimiter": ",",
    "skip_header": 1
}).csv("@AZURE_STAGE/Retail_Sales.csv")

df.show()

```

"c1"	"c2"	"c3"	"c4"	"c5"	"c6"	"c7"	"c8"	"c9"	"c10"	"c11"	"c12"	"c13"
ORD-SF8DGF0C	2024-10-08	2024-10	CUST1000	Ananya Sharma	India	South	Mumbai	Office Supplies	Paper	9	0.00	2700.
0	780.43											
ORD-BF0078E4	2024-08-11	2024-08	CUST1001	Aarav Iyer	India	Central	Lucknow	Technology	Networking	4	0.15	2720
0.0	4135.60											
ORD-86CD58A3	2024-06-12	2024-06	CUST1002	Arjun Sharma	USA	East	Kolkata	Furniture	Tables	4	0.10	3150
0.0	5676.96											
ORD-FR0C2D9	2024-12-18	2024-12	CUST1003	Ananya Das	India	North	Kolkata	Office Supplies	Appliances	9	0.00	3600
0.0	11783.22											
ORD-EF35506B	2024-10-27	2024-10	CUST1004	Ishaan Bhat	UK	Central	Chennai	Furniture	Storage	4	0.00	2400
0.0	1444.00											

- Renamed and casted columns to proper data types (OrderDate, Sales, Profit, etc.).
- Added a derived column **ProfitMargin** for analysis.
- Saved the final modeled table as **RETAIL_SALES_MODELED** in RETAIL_SCHEMA.

```

col("c12").cast("FLOAT").alias("Discount"),
col("c13").cast("FLOAT").alias("Sales"),
col("c14").cast("FLOAT").alias("Profit")
)

df_model.write.mode("overwrite").save_as_table("RETAIL_SCHEMA.RETAIL_SALES_MODELED")
print("Modeled data saved to RETAIL_SCHEMA.RETAIL_SALES_MODELED")

```

Modeled data saved to RETAIL_SCHEMA.RETAIL_SALES_MODELED

```

from snowflake.snowpark.functions import when

df_transformed = df_model.with_column(
    "ProfitMargin",
    when(col("Sales") != 0, (col("Profit") / col("Sales")) * 100).otherwise(0)
)

df_transformed.write.mode("overwrite").save_as_table("RETAIL_SCHEMA.RETAIL_SALES_MODELED")
print("Added ProfitMargin column and updated table")

```

Added ProfitMargin column and updated table

Step 3: Model and Structure Data in Snowflake

- Ensured the table **RETAIL_SALES_MODELED** contains clean, well-typed data.
- Verified column names, data types, and row counts.
- Ensured all key columns are correctly formatted (dates, floats, numbers).
- Confirmed transformations and calculations (e.g., Profit Margin).

The screenshot shows the Snowflake SQL Editor interface. At the top, the user is logged in as ACCOUNTADMIN and is working on a query named RETAIL_WH (X-Small). The query is a SELECT statement from the RETAIL_SALES_RAW table, filtered by a date range. The results are displayed in a table with 12 rows and 9 columns. The columns are ORDERID, ORDERDATE, MONTHOFSALE, CUSTOMERID, CUSTOMERNAME, COUNTRY, REGION, and C. The results show a list of orders with their details. On the right side, there is a 'Query Details' panel showing the query duration (2.4s), the number of rows (25), and the query ID (01bfe2a4-0001-673e-0...). Below this, there are two charts: a bar chart for ORDERID and a line chart for ORDERDATE.

	ORDERID	ORDERDATE	MONTHOFSALE	CUSTOMERID	CUSTOMERNAME	COUNTRY	REGION	C
1	ORD-5F8D6F0C	2024-10-08	2024-10	CUST1000	Ananya Sharma	India	South	Mum
2	ORD-BF0078E4	2024-08-11	2024-08	CUST1001	Aarav Iyer	India	Central	Luck
3	ORD-B6CD58A3	2024-06-12	2024-06	CUST1002	Arjun Sharma	USA	East	Kolka
4	ORD-FB0CD2D9	2024-12-18	2024-12	CUST1003	Ananya Das	India	North	Kolka
5	ORD-EF355968	2024-10-27	2024-10	CUST1004	Ishaan Bhat	UK	Central	Chen
6	ORD-60D1DA88	2024-08-26	2024-08	CUST1005	Neha Iyer	UAE	West	Chen
7	ORD-A5081404	2025-09-15	2025-09	CUST1006	Arjun Iyer	India	Central	Jaipu
8	ORD-E1C9BE42	2024-02-27	2024-02	CUST1007	Priya Singh	India	North	Luck
9	ORD-4FCB3B05	2025-05-26	2025-05	CUST1008	Kabir Menon	India	West	Jaipu
10	ORD-921966C8	2025-03-14	2025-03	CUST1009	Arjun Chopra	UAE	West	Mum
11	ORD-E4A002F0	2024-06-12	2024-06	CUST1010	Ananya Patel	UK	West	Ahme
12	ORD-0944D71F	2025-07-20	2025-07	CUST1011	Sanjay Gupta	UAE	North	Pune

Step 4: Build Power BI Report

- Connected Power BI Desktop to Snowflake using database and schema credentials.

The screenshot shows the 'Add data to your report' dialog box in Power BI Desktop. The dialog is titled 'Snowflake'. It has two input fields: 'Server' and 'Warehouse'. The 'Server' field contains the text 'TPRTVOG-TG33465.snowflakecomputing.com'. The 'Warehouse' field contains the text 'RETAIL_WH'. Below these fields, there is a link for 'Advanced options'. At the bottom right, there are two buttons: 'OK' and 'Cancel'.

Snowflake

Microsoft Account

KeyPair

Snowflake

tprtvog-tg33465.snowflakecomputing.com;RETAIL_...

User name

shivashankari

Password

●●●●●●●●●●●●●●●●

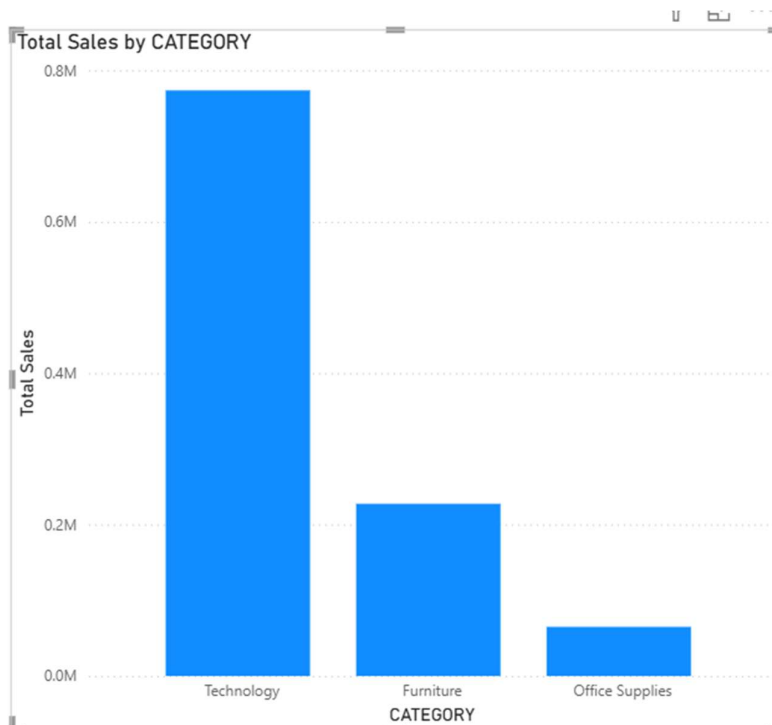
Back

Connect

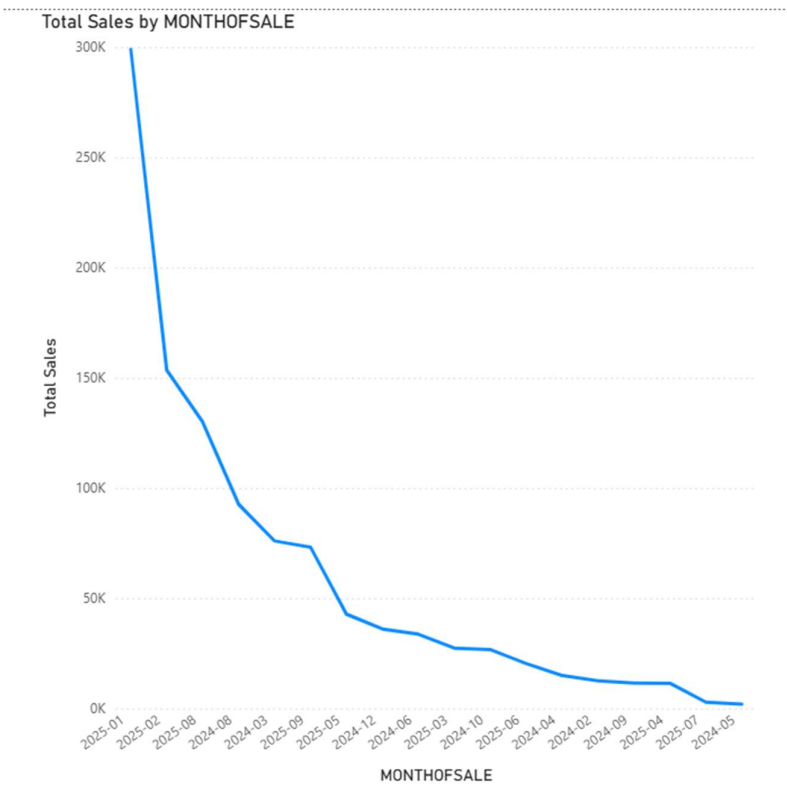
Cancel

- Selected table **RETAIL_SALES_MODELED** for data import.
- Built visuals using these measures:

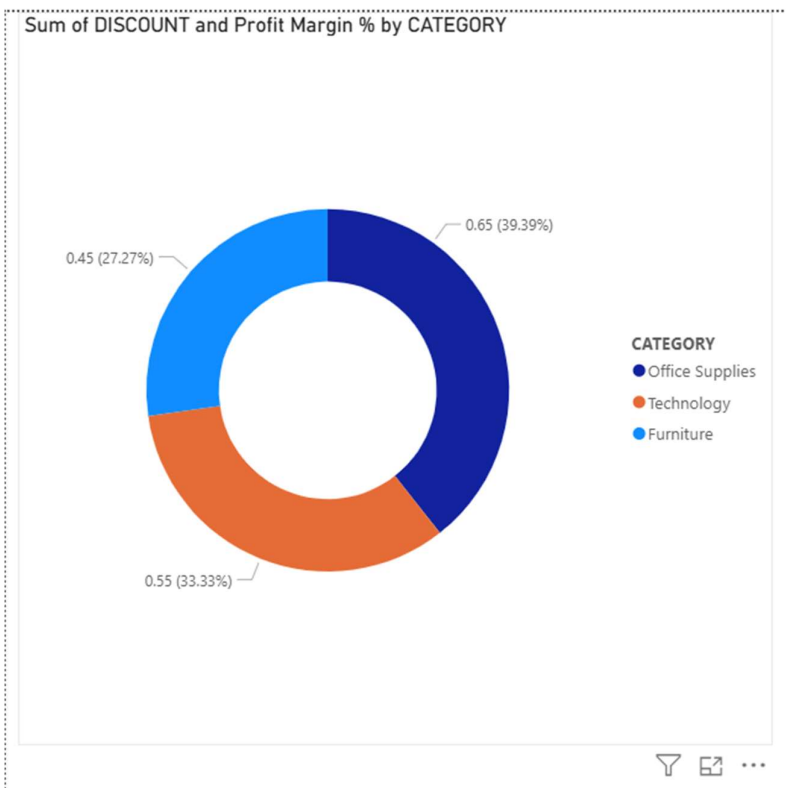
Column Chart: Sales by Category



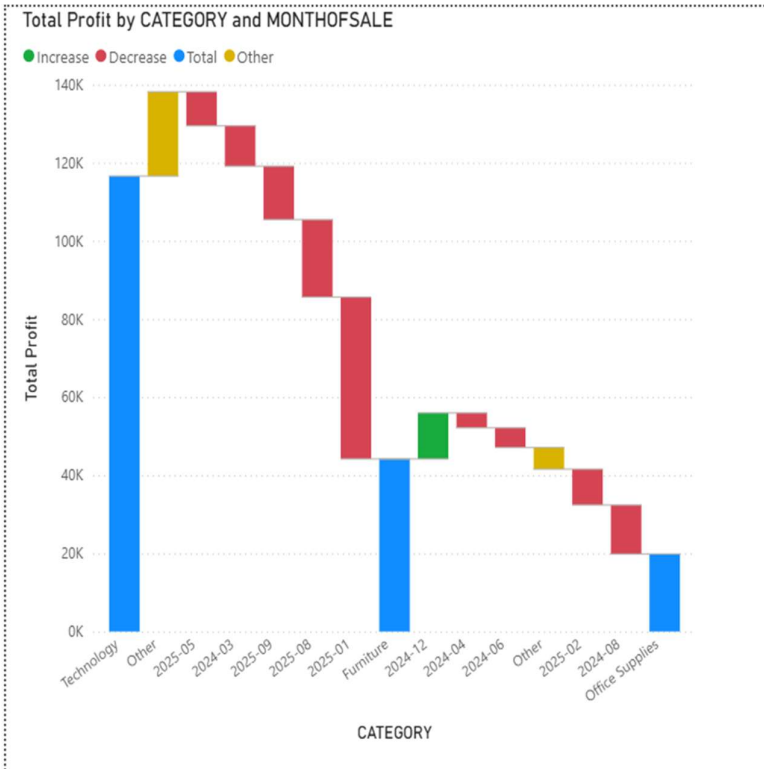
Line Chart: Monthly Sales Trend



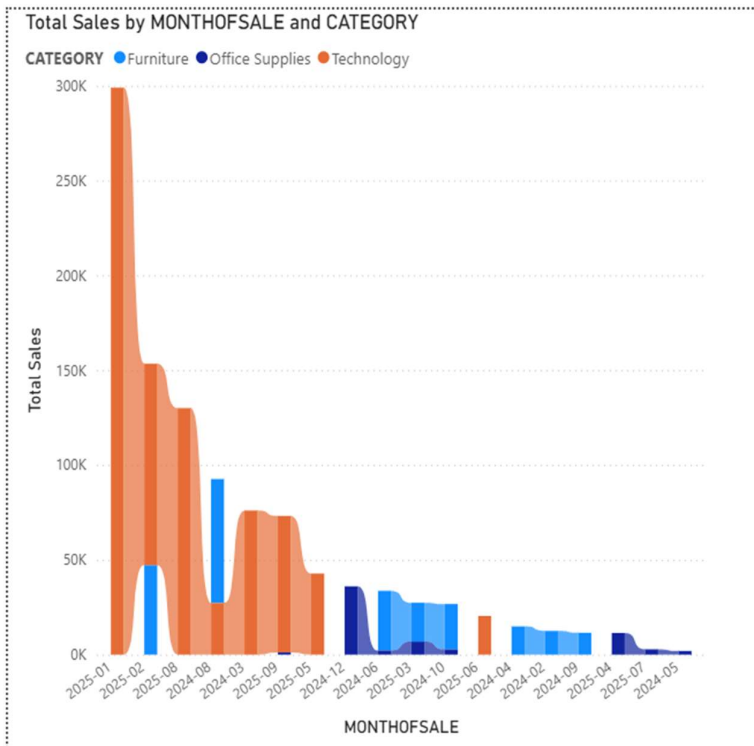
Donut Chart: Discount Impact



Waterfall Chart: Profit Breakdown



Ribbon Chart: Category rank changes over time



Area Chart : Show cumulative growth in sales or profit over time.

